



POLITECNICO
MILANO 1863

Protecting GPU applications via CUSPIS

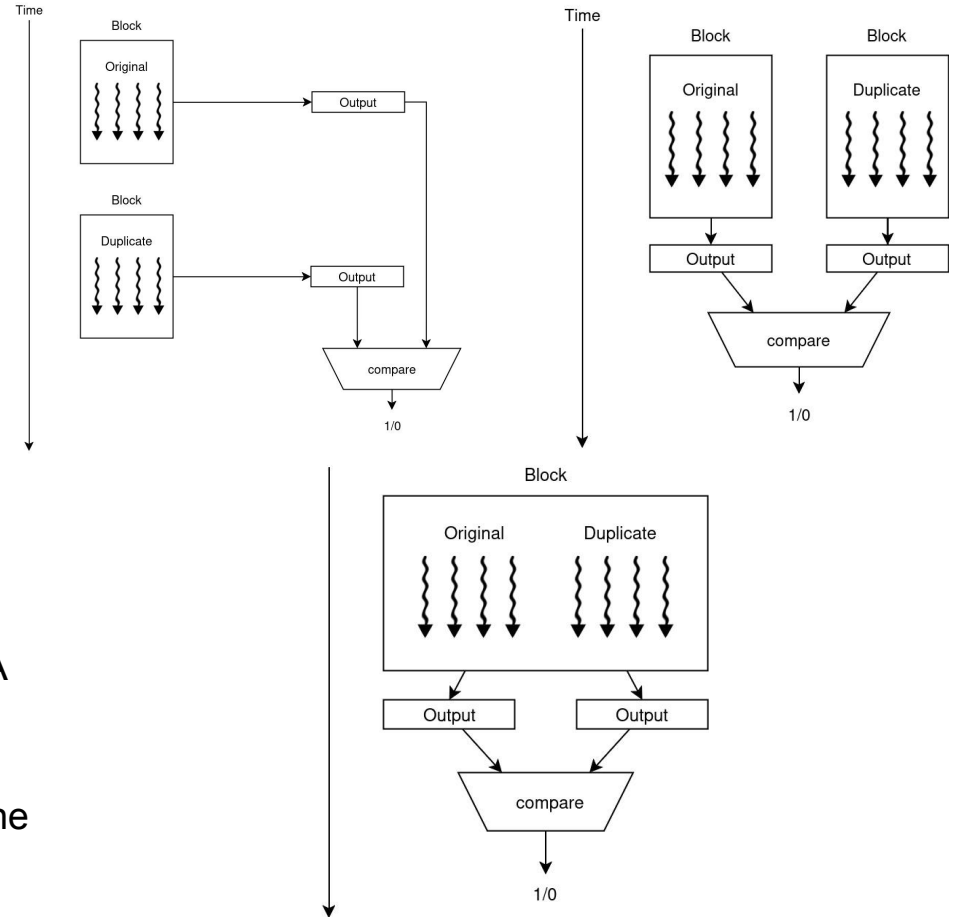
Francesco Galbiati 10728028,
Tiya Jreige 10958178



POLITECNICO
MILANO 1863

Context: CUSPIS

- Library that implements execution redundancy for CUDA kernels
- Makes available 3 kinds of policies:
 - Redundant Kernels
 - Redundant Blocks
 - Redundant Threads
- Memory allocations and kernel launch handled with dedicated wrappers of CUDA functions
- Redundant executions are compared by the library




Project Objectives

- Goal:
 - Implement redundancy mechanisms on a CUDA-based molecular dynamics benchmark.
- Benchmark
 - Simulates particle interaction in a 3D space
 - Divide the space into boxes
- Contribution:
 - Fault-tolerant GPU computations using CUSPIS library
 - Support of duplicated and triplicated executions


Problems

- The main problems encountered during the project were of two natures:
 - Problems regarding memory addresses:
 - Shared memory in thread redundancy was not enough for more than one replica
 - Indexes read in the replicated area of memory pointed to non-replicated memory locations
 - Several variables were accessed by thread or block index, generating conflicts in the case of replicas
 - Incompatibility between redundancy policies:
 - The different redundancy policies required different behaviours in some lines of code to correctly replicate the execution


Solution



Thread Redundancy



Block Redundancy



Solution

- Offset to access the first replicated particle with respect to the block/thread replica

```
- int block_replica = bx / d_dim_gpu.number_boxes + tx /  
  NUMBER_THREADS;
```

- Access to the appropriate box for the thread replica

```
- int thread_replica = tx / NUMBER_THREADS;
```

- Pointer to the first particle in a box:

```
- first_i = d_box_gpu[bx + thread_replica *  
  d_dim_gpu.number_boxes].offset + ( block_replica * NUMBER_PAR_PER_BOX *  
  d_dim_gpu.number_boxes);  
  
- first_j = d_box_gpu[pointer + thread_replica *  
  d_dim_gpu.number_boxes].offset + ( block_replica * NUMBER_PAR_PER_BOX *  
  d_dim_gpu.number_boxes);
```

Solution

- Modulo ensures that replicated threads start at an offset of 0 from the first particle in their box

```
- wtx = wtx % NUMBER_THREADS;
  while (wtx < NUMBER_PAR_PER_BOX) {
      rA_shared[wtx + NUMBER_PAR_PER_BOX * thread_replica] = rA[wtx];
      wtx = wtx + NUMBER_THREADS;
  }
  wtx = tx;
```

```
- wtx = wtx % NUMBER_THREADS;
  while (wtx < NUMBER_PAR_PER_BOX) {
      rB_shared[wtx + NUMBER_PAR_PER_BOX * thread_replica] = rB[wtx];
      // Copying data relates to particles
      wtx = wtx + NUMBER_THREADS;
  }
  wtx = tx;
```

Solution

- The “j” index was updated to access the appropriate neighbour replicas

```
wtx = wtx % NUMBER_THREADS;
while (wtx < NUMBER_PAR_PER_BOX) {
    for (j = NUMBER_PAR_PER_BOX * thread_replica; j <
        NUMBER_PAR_PER_BOX + (thread_replica * NUMBER_PAR_PER_BOX); j++)
    {
        // Shared memory is accessed with an index of
        // wtx + NUMBER_PAR_PER_BOX * thread_replica
        // for performing computations
    }
}
```


Concrete Outcomes

The concrete outcomes can be found in this GitHub [repository](#):

- The final product is under the name of `cuspis_lava_kbtr.cu`
- The repository also contains other files necessary to run the benchmark as well as the original CUDA version of it
- The project can be run either locally or on Google Collab using the Tesla T4 runtime

Learning outcomes:

- More familiar with CUDA and with GPU programming in general
- Explored an application of the concept of redundancy which was not part of the AOS or ES courses
- Soft skills relative to team work (work balance, communication and discussion skills)

Conclusions

- We managed to port the requested benchmark in CUSPIS and to obtain successful and correct executions
- Some limitations include further memory optimizations of some edge cases that occur according to the policies to be used
- Ultimately, we believe the final product of our work is ready to be used in a fault injection context